


MMLU compression efficiency

Retained (%)

x = GB saved, y = accuracy retained vs raw baseline



Models

- gemma-2-9b
- llama-3-8b
- mistral-7b
- phi-3

Quantizers

- i8-channel
- i8-tensor
- i4-grouped